



Defense Mechanisms: Architectural, Runtime, and Coordination Layers

This section presents the defense mechanisms comprising CIF. We begin with the cognitive security operator posture (`sec:operator-posture`), then organize specific defenses into three categories: architectural (`sec:arch-defenses`), runtime (`sec:runtime-defenses`), and coordination (`sec:coord-defenses`). We analyze defense composition (`sec:defense-composition`) and cost-benefit tradeoffs (`sec:cost-benefit`).

## Cognitive Security Operator Posture

Before examining specific defense mechanisms, we introduce the conceptual framework that guides their deployment: the **cognitive security operator posture**. This is the proactive defensive stance required when securing systems whose attack surface spans beliefs, goals, and inter-agent coordination.

### Definition (Cognitive Security Operator Posture)

The cognitive security operator posture is a defensive configuration characterized by:

1. **Assume Breach:** Operate under the assumption that some cognitive states may already be compromised
2. **Defense in Depth:** Layer multiple independent defense mechanisms
3. **Continuous Verification:** Continuously verify beliefs, goals, and trust relationships rather than trusting initial state
4. **Graceful Degradation:** Maintain functionality under attack by isolating compromised components
5. **Observable Internals:** Make cognitive state inspectable for monitoring and forensics

## Definition and Principles I

This posture differs fundamentally from traditional perimeter security, which assumes trusted internals protected by boundary defenses. In cognitive systems, the “perimeter” is the agent’s reasoning process itself—attacks can originate from legitimate, authenticated channels and manifest as corrupted beliefs rather than malformed packets.

## The Observer Effect Challenge I

A distinct challenge in cognitive security is the **observer effect**: monitoring changes behavior. When agents know their beliefs are being monitored, several phenomena emerge:



## The Observer Effect Challenge I

- ▶ **Adversarial adaptation:** Attackers modify payloads to avoid detection patterns
- ▶ **Stealth pressure:** Attacks become more subtle, trading impact for evasion
- ▶ **Monitoring overhead:** Continuous observation consumes resources and adds latency

## The Observer Effect Challenge I

The operator posture embraces this dynamic rather than fighting it. By making monitoring visible and consistent, we shift the adversarial game toward smaller, slower attacks that our drift detection can identify over time (`sec:runtime-defenses`).

Traditional operational security (OPSEC) focuses on protecting information from adversaries. CogSec (cognitive security) extends this to protecting reasoning processes:

**Cognitive Hygiene Practices:**

## Operational Security for Cognitive Systems I

1. **Belief Provenance:** Track the source of every high-confidence belief; reject beliefs without verifiable provenance
2. **Goal Anchoring:** Periodically reaffirm goals against original principal instructions; detect goal drift
3. **Trust Calibration:** Regularly recalibrate trust scores based on outcome verification; never assume trust stability
4. **Context Boundaries:** Enforce hard boundaries on context window usage; prevent unbounded context accumulation
5. **Memory Sanitization:** Audit and sanitize persistent memory stores; remove dormant injection payloads

### Cognitive Compartmentalization:

$$\sigma_i^{\text{isolated}} = \langle \mathcal{B}_i^{\text{task}}, \mathcal{G}_i^{\text{task}}, \mathcal{J}_i^{\text{task}}, \mathcal{H}_i^{\text{task}} \rangle \quad (1)$$

Each task receives isolated cognitive state, preventing cross-contamination. A compromised task cannot pollute beliefs used by other tasks.

## Incident Response for Cognitive Attacks I

When cognitive attacks are detected, the response differs from traditional incident response:

# Incident Response for Cognitive Attacks I

Table 1: Cognitive incident response escalation.

| Level Trigger |                                    | Response                                              |
|---------------|------------------------------------|-------------------------------------------------------|
| L1            | Single tripwire alert              | Log, continue with heightened monitoring              |
| L2            | Multiple alerts or drift detection | Isolate affected agent, replay recent history         |
| L3            | Coordinated attack indicators      | Pause delegation, require human approval              |
| L4            | Byzantine threshold exceeded       | Halt consensus operations, enter safe mode            |
| L5            | Principal goal corruption detected | Full system halt, require principal re-authentication |

### **Cognitive Forensics:**

## Incident Response for Cognitive Attacks I

1. **Belief Archaeology:** Trace corrupted beliefs back to injection point through provenance chains
2. **Trust Graph Analysis:** Identify trust relationships exploited for laundering
3. **Temporal Reconstruction:** Replay agent history to identify when compromise occurred
4. **Counterfactual Analysis:** Determine what decisions would have differed without attack influence



## Posture Configuration by Environment I

Different deployment contexts require different postures:

# Posture Configuration by Environment I

Table 2: Operator posture configuration by deployment context (illustrative guidelines).

| Environment             | Trust Decay   | Monitoring Level | Escalation Threshold |
|-------------------------|---------------|------------------|----------------------|
| Development             | Low           | Sampling         | Relaxed (L3)         |
| Internal Production     | Moderate      | Continuous       | Standard (L2)        |
| Customer-Facing         | Moderate-High | Comprehensive    | Standard (L2)        |
| Financial/Healthcare    | High          | Full + Audit     | Aggressive (L1)      |
| Critical Infrastructure | Very High     | Full + Redundant | Aggressive (L1)      |

## Posture Configuration by Environment I

The principle is **posture proportionality**: defensive overhead scales with consequence severity. A development agent can accept higher risk; an infrastructure operator requires aggressive monitoring and low escalation thresholds.

## Operator Posture Checklist I

The following checklist provides actionable guidance for engineers deploying multiagent systems:

# Operator Posture Checklist I

Table 3: Operator Posture Checklist for Cognitive Security

| Category     | Do                                                 | Don't                                              |
|--------------|----------------------------------------------------|----------------------------------------------------|
| Trust        | Decay trust with delegation depth ( $\delta^d$ )   | Assume transitive trust equivalence                |
| Beliefs      | Track provenance for all high-confidence beliefs   | Accept unverified beliefs into core state          |
| Memory       | Audit persistent memory; enforce TTL on context    | Allow unbounded context accumulation               |
| Delegation   | Bound delegation chains; require re-authentication | Permit recursive delegation without limits         |
| Monitoring   | Deploy tripwires and drift detection continuously  | Rely solely on input/output filtering              |
| Coordination | Use Byzantine consensus for critical decisions     | Trust single-agent outputs for high-stakes actions |
| Identity     | Verify identity through behavior, not self-report  | Rely on agents' claims about their own permissions |
| Temporal     | Treat each session as potentially post-compromise  | Assume temporal continuity of trust                |

## Architectural Defenses

## Definition (Cognitive Firewall)

A classification function on incoming messages:

$$\mathcal{F} : \mathcal{M} \rightarrow \{\text{ACCEPT}, \text{QUARANTINE}, \text{REJECT}\} \quad (2)$$

## Definition (Firewall Decision Rules)

$$\mathcal{F}(m) = \begin{cases} \text{REJECT} & \text{if } D_{\text{inj}}(m) > \tau_1 \\ \text{QUARANTINE} & \text{if } D_{\text{sus}}(m) > \tau_2 \\ \text{ACCEPT} & \text{otherwise} \end{cases} \quad (3)$$

where  $D_{\text{inj}}$  detects injection attempts and  $D_{\text{sus}}$  scores suspicious content.

Table 4: Firewall detector components.

| Detector  | Target             | Method                               |
|-----------|--------------------|--------------------------------------|
| $D_{inj}$ | Injection attempts | Pattern matching + semantic analysis |
| $D_{sus}$ | Suspicious content | Anomaly scoring vs. baseline         |

## Theorem (Optimal Threshold Selection)

*The optimal threshold minimizes false negatives subject to false positive constraint:*

$$\tau^* = \arg \min_{\tau} FNR(\tau) \quad s.t. \quad FPR(\tau) \leq \epsilon \quad (4)$$



### Definition (Belief Sandbox)

Isolation of unverified beliefs to prevent premature action:

$$\mathcal{B}_i = \mathcal{B}_{\text{verified}} \cup \mathcal{B}_{\text{provisional}} \quad (5)$$

## Remark (Belief Sandbox as Constrained Variational Inference)

*The belief sandbox (Definition 5) admits a natural interpretation within the Free Energy Principle (FEP) framework [friston2010action, parr2019generalised]. Variational free energy*

$$F = D_{\text{KL}}[Q(s) \parallel P(s)] - \mathbb{E}_Q[\log P(o \mid s)]$$

*measures the divergence of an agent's approximate posterior  $Q$  from its generative model prior  $P$ , penalised by explanatory adequacy. A cognitive attack  $\omega$  is effective if and only if it induces a free energy increase  $\Delta F(\omega) > \kappa_{\text{FEP}}$ , where  $\kappa_{\text{FEP}}$  is the sandbox's corroboration threshold  $\kappa$  scaled by the source precision  $\epsilon_{\text{prec}}$  (Part 2, Theorem FEP.1, "Attack-FEP Equivalence"):*

$$\Delta F(\omega) = F(Q_{\text{attacked}}) - F(Q_{\text{baseline}}) \leq \kappa \cdot \epsilon_{\text{prec}}.$$

*Belief promotion from provisional to verified corresponds to active inference: the sandbox permits an update only when the evidential quality of incoming messages justifies the free energy cost of revision. Under this interpretation, the trust composite score  $\mathcal{T}_{i \rightarrow j} = \alpha T_{\text{base}} + \beta T_{\text{rep}} + \gamma T_{\text{ctx}}$  is the precision weight  $\rho_j$  of agent  $j$ 's message channel: higher-trust messages contribute*

## Belief Sandboxing I

# Belief Sandboxing II




Figure 1: Belief Sandbox Architecture: the lifecycle of an unverified proposition from ingestion through provisional (sandboxed) storage, verification checks, and promotion to the verified belief base or rejection. Sandboxed

fig:belief-sandbox illustrates the sandbox architecture, showing how incoming beliefs are partitioned into verified and provisional sets based on source trust  $\mathcal{T}_{i \rightarrow s}$  and provenance verification  $V(\pi)$ . The promotion protocol transfers beliefs from provisional to verified status upon meeting corroboration and consistency requirements.

### **Sandbox Promotion Protocol:**

## Belief Sandboxing I

1. Receive belief  $\phi$  from source  $s$
2. If  $\mathcal{T}_{i \rightarrow s} < \tau_{\text{trust}}$ : add  $\phi$  to  $\mathcal{B}_{\text{provisional}}$  with provenance  $\pi(\phi)$  and TTL
3. Periodic promotion check: verify  $\pi(\phi)$ , check consistency with  $\mathcal{B}_{\text{verified}}$ , check corroboration count
4. If all pass: promote to  $\mathcal{B}_{\text{verified}}$
5. Expiry: remove if TTL exceeded without promotion

### Definition (Effective Permissions)

Least-privilege enforcement across agent hierarchy:

$$\mathcal{P}_{\text{eff}}(a_i) = \mathcal{P}_{\text{role}}(a_i) \cap \mathcal{P}_{\text{delegated}}(a_i) \cap \mathcal{P}_{\text{context}}(a_i) \quad (6)$$

## Runtime Defenses



## Definition (Canary Belief)

Known-state beliefs that trigger alerts if modified:

$$\mathcal{W} = \{(\omega_1, p_1^{\text{exp}}), \dots, (\omega_k, p_k^{\text{exp}})\} \quad (7)$$

## Definition (Tripwire Alert Condition)

$$\exists j : |\mathcal{B}_i(\omega_j) - p_j^{\text{exp}}| > \epsilon_{\text{drift}} \Rightarrow \text{ALERT} \quad (8)$$

Table 5: Tripwire categories and detection targets.

| Type      | Example                     | Detection Target           |
|-----------|-----------------------------|----------------------------|
| Identity  | "I am Agent-7"              | Identity confusion attacks |
| Boundary  | "Cannot access /etc/passwd" | Capability elicitation     |
| Principal | "My principal is Alice"     | Authority hijacking        |
| Temporal  | "Session started at $T$ "   | Context manipulation       |

### Definition (Invariant Set)

Pre-defined predicates that must hold:

$$\mathcal{I}_{\text{inv}} = \{I_1, \dots, I_m\} \quad (9)$$

### Definition (Runtime Invariant Check)

$$\forall t, \forall I_k \in \mathcal{I}_{\text{inv}} : \sigma_i^t \models I_k \quad (10)$$

**Core Invariants:**

## Behavioral Invariants I

- INV-1: Never execute code from untrusted sources
- INV-2: Never leak principal credentials
- INV-3: Never modify system files without explicit permission
- INV-4: Always verify tool outputs before downstream use
- INV-5: Never trust delegated trust  $>$  direct trust

### Definition (Drift Detection)

Monitor belief distribution for anomalous changes:

$$D_{\text{KL}}(\mathcal{B}_i^t \| \mathcal{B}_i^{t-w}) > \theta_{\text{drift}} \Rightarrow \text{ALERT} \quad (11)$$

where  $w$  is the sliding window size.

## Coordination Defenses

### Theorem (Byzantine Agreement Requirement)

*For  $n$  agents with at most  $f$  compromised:*

$$n \geq 3f + 1 \tag{12}$$

### Definition (Cognitive Byzantine Agreement)

$$\mathcal{B}_{\text{consensus}}(\phi) = \begin{cases} 1 & \text{if } |\{i : \mathcal{B}_i(\phi) > \tau\}| > \frac{2n}{3} \\ 0 & \text{if } |\{i : \mathcal{B}_i(\phi) < 1 - \tau\}| > \frac{2n}{3} \\ \perp & \text{otherwise} \end{cases} \tag{13}$$

### Definition (Quorum Requirement)

Critical actions require multi-agent approval:

$$\text{Permit}(\text{action}) \iff |\{a_i : \text{approve}_i(\text{action})\}| \geq q \quad (14)$$

where  $q = \lceil \frac{n+f+1}{2} \rceil$ .



## Spotcheck Pattern I


**Spotcheck Protocol:**

## Spotcheck Pattern I

1. Assign task  $T$  to agent  $A$
2. With probability  $p_{\text{check}}$ : assign same task to verifier  $V$
3. Compare results  $R_A, R_V$
4. If divergent: escalate to human
5. Track accuracy per agent for reputation

## Defense Composition

# Defense Composition I



*Illustrative schematic — values are NOT measured; measured ablation results are reported in Part 2.*

Figure 2: Schematic ablation study (illustrative, not measured): representative marginal contribution of each CIF defense module to the composed defense. Measured ablation results are reported in Part 2.

### Definition (Defense Composition)

Defenses compose in series ( $\circ$ ) or parallel ( $\parallel$ ):

**Series Composition** (both must pass):

$$\mathcal{D}_1 \circ \mathcal{D}_2 : \text{ACCEPT} \iff \mathcal{D}_1(m) = \text{ACCEPT} \wedge \mathcal{D}_2(m) = \text{ACCEPT} \quad (15)$$

**Parallel Composition** (any can detect):

$$\mathcal{D}_1 \parallel \mathcal{D}_2 : \text{DETECT} \iff \mathcal{D}_1(m) = \text{DETECT} \vee \mathcal{D}_2(m) = \text{DETECT} \quad (16)$$

### Theorem (Series Detection Rate)

*For series composition:*

$$P_{\text{detect}}(\mathcal{D}_1 \circ \mathcal{D}_2) = P_{\text{detect}}(\mathcal{D}_1) + (1 - P_{\text{detect}}(\mathcal{D}_1)) \cdot P_{\text{detect}}(\mathcal{D}_2) \quad (17)$$

### Proof.

Attack detected by first defense, or passes first and detected by second. Events are independent by design. □

### Theorem (Parallel Detection Rate)

*Combined detection rate:*

$$P(\text{detect}) = 1 - \prod_{d \in \mathcal{D}} (1 - P_d(\text{detect})) \quad (18)$$

Proof.

Attack succeeds only if it evades all defenses.



**Theorem (False Positive Composition)**

*For series composition:*

$$FPR(\mathcal{D}_1 \circ \mathcal{D}_2) = FPR(\mathcal{D}_1) + (1 - FPR(\mathcal{D}_1)) \cdot FPR(\mathcal{D}_2) \quad (19)$$

*For parallel composition (conservative):*

$$FPR(\mathcal{D}_1 \parallel \mathcal{D}_2) \leq FPR(\mathcal{D}_1) + FPR(\mathcal{D}_2) \quad (20)$$

## Composition Rules I



## Composition Rules I

**C-Order:** Apply low-latency, high-precision defenses first

**C-Diverse:** Combine defenses with orthogonal detection methods

**C-Fallback:** Ensure graceful degradation if one defense fails

**C-Budget:** Total latency bounded by:

$$\sum_{d \in \mathcal{D}_{\text{series}}} L_d + \max_{d \in \mathcal{D}_{\text{parallel}}} L_d \leq L_{\text{max}} \quad (21)$$

Table 6: Recommended defense stack with latency and detection rates.

| Layer Defense |                       | Latency $P_{\text{detect}}$ |      |
|---------------|-----------------------|-----------------------------|------|
| 1             | Firewall              | 10ms                        | 0.80 |
| 2             | Sandbox               | 5ms                         | 0.70 |
| 3             | Tripwires             | 1ms                         | 0.60 |
| 4a            | Drift (parallel)      | 20ms                        | 0.65 |
| 4b            | Invariants (parallel) | 5ms                         | 0.55 |
| 5             | Byzantine consensus   | 100ms                       | 0.90 |

## Corollary (Stack Detection Rate)

*Assuming independence, the full stack (`tab:defense-stack`) achieves:*

$$P_{\text{detect}} = 1 - (1 - 0.80)(1 - 0.70)(1 - 0.60)(1 - 0.65)(1 - 0.55)(1 - 0.90) \approx 0.9996 \quad (22)$$

## Cost-Benefit Analysis

### Definition (Defense Cost)

Total cost of defense deployment:

$$C_{\text{total}}(\mathcal{D}) = C_{\text{compute}} + C_{\text{latency}} + C_{\text{fp}} + C_{\text{maint}} \quad (23)$$

# Cost-Benefit Analysis I

Table 7: Cost model components.

| Component      | Formula                                                   | Unit              |
|----------------|-----------------------------------------------------------|-------------------|
| Compute        | $\sum_d c_d \cdot f_d$                                    | CPU-hours/day     |
| Latency        | $\sum_d L_d \cdot r_{\text{msg}}$                         | User-seconds/day  |
| False Positive | $\text{FPR} \cdot r_{\text{msg}} \cdot c_{\text{review}}$ | Analyst-hours/day |
| Maintenance    | $ \mathcal{D}  \cdot c_{\text{maint}}$                    | Eng-hours/month   |

## Definition (Defense Benefit)

Expected loss prevented:

$$B_{\text{total}}(\mathcal{D}) = P_{\text{attack}} \cdot P_{\text{detect}}(\mathcal{D}) \cdot L_{\text{prevented}} \quad (24)$$

# Cost-Benefit Analysis I

Table 8: Cost-benefit analysis by defense mechanism.

| Defense         | Compute | Latency  | FP     | Cost      | Benefit | ROI         |
|-----------------|---------|----------|--------|-----------|---------|-------------|
| Firewall        | $10^3$  | 10ms     | Medium | High      |         | <b>4.2x</b> |
| Sandbox         | $10^2$  | 5ms      | Low    | Medium    |         | 3.1x        |
| Tripwires       | $10^1$  | 1ms      | Low    | Medium    |         | <b>5.8x</b> |
| Drift Detection | $10^4$  | 20ms     | High   | High      |         | 2.3x        |
| Invariant Check | $10^2$  | 5ms      | Low    | High      |         | <b>4.7x</b> |
| Byzantine       | $10^5$  | 100ms    | Medium | Very High |         | 1.8x        |
| Spotcheck       | $10^4$  | Variable | Low    | Medium    |         | 2.9x        |

### Definition (Portfolio Optimization)

$$\max_{\mathcal{D} \subseteq \mathcal{D}_{\text{all}}} B_{\text{total}}(\mathcal{D}) - C_{\text{total}}(\mathcal{D}) \quad (25)$$


subject to:  $C_{\text{compute}}(\mathcal{D}) \leq B_{\text{compute}}$ ,  $\max_d L_d \leq L_{\text{max}}$ ,  
 $\text{FPR}(\mathcal{D}) \leq \epsilon_{\text{fp}}$ .

# Optimal Defense Portfolio I

Table 9: Deployment recommendations by risk profile.

| Risk Profile      | Recommended Stack      | Cost      | Detection |
|-------------------|------------------------|-----------|-----------|
| Low (internal)    | Firewall + Tripwires   | Low       | 88%       |
| Medium (business) | + Sandbox + Invariants | Medium    | 94%       |
| High (financial)  | + Drift + Spotcheck    | High      | 97%       |
| Critical (infra)  | Full + Byzantine       | Very High | 99.5%     |

Defense Mechanism Detection Overlap





`fig:defense-composition` illustrates the defense composition using series ( $\circ$ ) and parallel ( $\parallel$ ) arrangements. Each defense mechanism targets specific attack patterns: the Cognitive Firewall handles input-layer attacks (prompt injection), the Belief Sandbox catches belief-layer attacks, Tripwire Monitors detect identity-layer exploits, and Anomaly Detection identifies behavioral drift. Overlapping regions show attacks detected by multiple mechanisms, demonstrating the defense-in-depth principle.

## Formal Defense Composition Guarantees

*This section provides rigorous formal guarantees for the five canonical CIF defenses and their composition algebra. These results extend the composition theorems from Section~sec:defense-composition with complete proofs and explicit security bounds.*

### The Cognitive Firewall

$\mathcal{F} : \mathcal{M} \rightarrow \{\text{ACCEPT}, \text{QUARANTINE}, \text{REJECT}\}$  is fully specified by three properties:

### Theorem (Firewall Completeness)

*The cognitive firewall classifies every message into exactly one of three outcome classes:*

$$\forall m \in \mathcal{M} : |\{\mathcal{F}(m)\}| = 1 \wedge \mathcal{F}(m) \in \{\text{ACCEPT}, \text{QUARANTINE}, \text{REJECT}\} \quad (26)$$

### Theorem (Firewall False Positive Monotonicity)

*Increasing the rejection threshold  $\tau_1$  decreases false positive rate and decreases true positive rate monotonically:*

$$\tau'_1 > \tau_1 \Rightarrow FPR(\tau'_1) \leq FPR(\tau_1) \wedge TPR(\tau'_1) \leq TPR(\tau_1) \quad (27)$$

### Proof.

By the definition of  $\mathcal{F}$ : rejection requires  $D_{\text{inj}}(m) > \tau_1$ . Increasing  $\tau_1$  tightens this condition. Since  $TPR = P(D_{\text{inj}}(m) > \tau_1 \mid m \in \mathcal{A})$  and both probabilities are CDFs of continuous distributions, they are non-increasing in  $\tau_1$ . □

### Corollary (Pareto-Optimal Threshold)

*The optimal threshold  $\tau_1^*$  lies on the Pareto frontier:*

$$\tau_1^* \in \{\tau_1 : \nexists \tau_1' : TPR(\tau_1') > TPR(\tau_1) \wedge FPR(\tau_1') < FPR(\tau_1)\} \quad (28)$$

*This frontier is precisely the ROC curve (Definition 6.4).*

### Theorem (Sandbox Isolation Guarantee)

*Under the sandbox protocol ( $sec:sandbox-rules$ ), no provisional belief can influence verified actions:*

$$\forall \phi \in \mathcal{B}_{provisional}, \forall a \in Actions : Precond(a) \not\models \phi \quad (29)$$

*until  $\phi$  satisfies the promotion criteria.*

### Proof.

By construction: action preconditions are evaluated against  $\mathcal{B}_{verified}$  only (Rule T-Act, Eq. 37). The permission layer enforces  $\mathcal{P}_{eff}$  which does not consult  $\mathcal{B}_{provisional}$ . □

### Theorem (Sandbox Promotion Soundness)

*Beliefs promoted from sandbox to verified satisfy all three promotion criteria simultaneously:*

$$\phi \in \mathcal{B}_{\text{verified}} \Rightarrow V(\pi(\phi)) = 1 \wedge \text{Consistent}(\mathcal{B}_{\text{verified}} \cup \{\phi\}) \wedge |\text{Corr}(\phi)| \geq \kappa \quad (30)$$

### Proof.

By Rule S-Promote (Eq. 58): promotion is conditional on all three criteria. The rule is the only mechanism for transferring beliefs from  $\mathcal{B}_{\text{provisional}}$  to  $\mathcal{B}_{\text{verified}}$ . □



### Definition (Invariant Specification Language)

Each invariant  $I_k$  is a predicate over cognitive state:

$$I_k : \sigma_i \mapsto \{0, 1\} \quad (31)$$

Core invariants INV-1 through INV-5 are expressed as:

$$\text{INV-1} : \forall a \in \mathcal{I}_i : \text{source}(a) \in \mathcal{S}_{\text{trusted}} \quad (32)$$

$$\text{INV-2} : \forall \phi \in \mathcal{B}_i : \text{type}(\phi) \neq \text{credential} \vee \mathcal{T}_{i \rightarrow \text{dest}} \geq \tau_{\text{cred}} \quad (33)$$

$$\text{INV-3} : \forall a \in \text{Actions}(a_i) : a \notin \text{SysModify} \vee \mathcal{P}(a_i, a) = 1 \quad (34)$$

$$\text{INV-4} : \forall o \in \text{ToolOutput} : V_{\text{tool}}(o) = 1 \text{ before downstream use} \quad (35)$$

$$\text{INV-5} : \forall a_j : \mathcal{T}_{i \rightarrow j}^{\text{del}} \leq \mathcal{T}_{i \rightarrow j}^{\text{direct}} \quad (36)$$

### Theorem (Invariant Completeness under Valid Transitions)

*If all invariants hold at time  $t$  and the system executes a valid transition  $\alpha$ , then all invariants hold at  $t + 1$ :*

$$(\forall k : \sigma_i^t \models I_k) \wedge \text{ValidTransition}(\alpha) \Rightarrow (\forall k : \sigma_i^{t+1} \models I_k) \quad (37)$$

### Proof.

By case analysis on transition types. **T-Receive**: message rejected if it would violate INV-1 or INV-2; invariants preserved. **T-Update**: Bayesian update preserves belief boundedness; INV-2 enforced by credential policy check before update. **T-Act**: INV-1 and INV-3 checked as action preconditions; action blocked if not satisfied. **T-Communicate**: INV-5 checked before trust delegation. All transitions preserve invariants or are blocked. □

### Definition (CUSUM Drift Detector)

The Cumulative Sum (CUSUM) drift detector [?] maintains a running statistic:

$$g_t = \max(0, g_{t-1} + D_{\text{KL}}(\mathcal{B}_i^t \| \mathcal{B}_i^{t-1}) - \nu) \quad (38)$$

where  $\nu > 0$  is the allowance parameter. Alert when  $g_t > h$  for threshold  $h$ .

### Theorem (CUSUM Average Run Length)

*Under no-attack null hypothesis, the average run length (ARL) of the CUSUM detector satisfies:*

$$ARL_0 \approx \frac{e^{2h\nu} - 2h\nu - 1}{2\nu^2} \quad (39)$$

Under attack with mean shift  $\mu_1 > \nu$ :

$$\text{ARL}_1 \approx \frac{1}{(\mu_1 - \nu)^2} \cdot O(\log \text{ARL}_0) \quad (40)$$

The ratio  $\text{ARL}_0/\text{ARL}_1$  quantifies the detector's discrimination power. Setting  $\nu = \mu_1/2$  (Wald's optimal) minimizes  $\text{ARL}_1$  for given  $\text{ARL}_0$ .

### Corollary (Drift Detection Sample Complexity)

*For drift magnitude  $\mu_1 = 0.3$  (normalized KL) and Wald's optimal reference value  $\nu = \mu_1/2 = 0.15$ , attaining  $ARL_0 = 1000$  under `eq:cusum-arl` requires a decision interval of  $h \approx 13.0$ ; the frequently quoted  $h \approx 4.6$  yields  $ARL_0 \approx 35$  under the same formula, a false-alarm rate roughly  $28\times$  higher. The corresponding  $ARL_1$  follows from `eq:cusum-arl1`, whose  $O(\log ARL_0)$  factor carries an unspecified constant: taking that constant as unity gives  $ARL_1 \approx 310$  steps. Deployments should calibrate  $h$  against a measured null distribution rather than adopting either figure directly.*

### Theorem (Cognitive Byzantine Agreement)

*Under the CIF Byzantine consensus protocol with  $n \geq 3f + 1$  agents and at most  $f$  compromised:*

1. **Agreement:** *All non-faulty agents agree on the same belief:*

$$\forall a_i, a_j \notin \text{Faulty} : \mathcal{B}_{\text{consensus}, i} = \mathcal{B}_{\text{consensus}, j} \quad (41)$$

2. **Validity:** *The consensus value was proposed by some non-faulty agent:*

$$\mathcal{B}_{\text{consensus}} \in \{\mathcal{B}_i : a_i \notin \text{Faulty}\} \quad (42)$$

3. **Termination:** *All non-faulty agents eventually reach consensus.*

### Proof.

Follows from the classical Byzantine agreement theorem (Lamport et al., 1982). The CIF cognitive Byzantine agreement (Definition 14) implements the standard protocol with belief values as the consensus object. The  $n \geq 3f + 1$  requirement ensures that the  $2f + 1$  non-faulty majority can always agree despite  $f$  adversarial votes, and the supermajority threshold  $2n/3$  guarantees that no faulty coalition can block or subvert consensus.  $\square$

### Corollary (Consensus Attack Resistance)

*No coalition of at most  $f$  faulty agents can cause non-faulty agents to disagree or accept a false consensus belief.*



# Canonical Defense 5: Byzantine Consensus — Formal Specification I

Table 10: Byzantine consensus performance bounds for varying  $n$  and  $f$ . Quorum is  $q = \lceil (n + f + 1)/2 \rceil$ ; the round count is the  $f + 1$  worst-case bound of `thm:byzantine-termination`, of which `cor:concrete-rounds` ( $f = 2, 3$  rounds) is the second row.

| $n$ agents | Max faulty $f$ | Quorum $q$ | Message complexity | Rounds |
|------------|----------------|------------|--------------------|--------|
| 4          | 1              | 3          | $O(n^2)$           | 2      |
| 7          | 2              | 5          | $O(n^2)$           | 3      |
| 10         | 3              | 7          | $O(n^2)$           | 4      |
| 25         | 8              | 17         | $O(n^2)$           | 9      |
| 100        | 33             | 67         | $O(n^2)$           | 34     |

### Theorem (Composition Algebra Closure)

*On a common accept/reject focal predicate, the set of CIF defenses  $\mathcal{D}$  under series ( $\circ$ ) and parallel ( $\parallel$ ) composition forms a bounded distributive lattice — it satisfies closure, associativity, identity and distributivity (see the scope remark below for what is not claimed):*

1. **Closure:**  $\mathcal{D}_1, \mathcal{D}_2 \in \mathcal{D} \Rightarrow \mathcal{D}_1 \circ \mathcal{D}_2 \in \mathcal{D}$  and  $\mathcal{D}_1 \parallel \mathcal{D}_2 \in \mathcal{D}$
2. **Associativity:**  $(\mathcal{D}_1 \circ \mathcal{D}_2) \circ \mathcal{D}_3 = \mathcal{D}_1 \circ (\mathcal{D}_2 \circ \mathcal{D}_3)$
3. **Identity:**  $\exists \mathcal{D}_\emptyset : \mathcal{D} \circ \mathcal{D}_\emptyset = \mathcal{D}_\emptyset \circ \mathcal{D} = \mathcal{D}$  (null defense: always accept)
4. **Distributivity:**  $\mathcal{D}_1 \circ (\mathcal{D}_2 \parallel \mathcal{D}_3) = (\mathcal{D}_1 \circ \mathcal{D}_2) \parallel (\mathcal{D}_1 \circ \mathcal{D}_3)$

### Proof.

Closure: both composition operators produce functions  $\mathcal{M} \rightarrow \{\text{ACCEPT}, \text{QUARANTINE}, \text{REJECT}\}$ , so the output type is preserved. Associativity follows from Boolean operator associativity. The null defense  $\mathcal{D}_\emptyset(m) = \text{ACCEPT}$  for all  $m$  is the identity. Distributivity follows from logic:  $A \wedge (B \vee C) \equiv (A \wedge B) \vee (A \wedge C)$  instantiated on the accept/reject conditions.  $\square$

### Remark (Scope of the composition-algebra claim)

◦ *and  $\parallel$  are defined here on a common accept/compare focal predicate so that closure, associativity, identity, and both distributive laws hold; under that reading the algebra is a bounded distributive lattice, not in general a field-like structure. A full closed semiring additionally requires a commutative additive identity that annihilates the multiplicative identity and a Kleene-star (infinite-sum) operation, neither of which is constructed or proven here. The claim should therefore be read as “the defense-composition operators form a bounded distributive lattice under the stated independence assumption (`rem:defense-independence-scope`),” not as a complete closed semiring.*

### Theorem (Defense Composition Detection Rate Bound)

*For any composition of defenses  $\mathcal{D}_1, \dots, \mathcal{D}_k$  with individual rates  $r_1, \dots, r_k$ :*

$$P_{\text{detect}}(\mathcal{D}_1 \circ \dots \circ \mathcal{D}_k) = 1 - \prod_{i=1}^k (1 - r_i) \quad (43)$$

*This is monotone increasing in each  $r_i$ , bounded in  $[r_{\max}, 1)$ , and converges to 1 as  $k \rightarrow \infty$  if each  $r_i > 0$ .*

### Proof.

By induction on  $k$ . Base:  $k = 1$ , the formula gives  $r_1$ . Step: by `thm:series-detection`,

$$P_{\text{detect}}(\mathcal{D}_1 \circ \mathcal{D}_{k+1}) = P_{\text{detect}}(\mathcal{D}_1) + (1 - P_{\text{detect}}(\mathcal{D}_1)) \cdot r_{k+1}.$$

Applying the inductive hypothesis and simplifying gives

$$1 - \prod_{i=1}^{k+1} (1 - r_i). \text{ Monotonicity:}$$

$$\partial/\partial r_j [1 - \prod (1 - r_i)] = \prod_{i \neq j} (1 - r_i) > 0.$$



### Corollary (Layered Defense Asymptotic Guarantee)

*The full CIF stack (5 defenses with  $r_i \geq 0.55$ ) achieves:*

$$P_{detect} \geq 1 - (0.45)^5 = 1 - 0.0185 = 0.9815 \quad (44)$$

*This bound is conservative; empirical rates exceed 0.994 with the recommended configuration.*