

Formal Verification: Safety Properties and Model Checking

This section establishes safety properties
(`sec:safety-properties`), proves invariant preservation lemmas
(`sec:invariant-lemmas`), demonstrates liveness guarantees
(`sec:liveness-properties`), derives complexity bounds
(`sec:complexity-bounds`), and presents model checking
verification (`sec:model-checking`).

Safety Properties

Theorem (Belief Injection Resistance)

Under CIF with firewall detection rate r_f and sandboxing verification rate r_s :

$$P(\mathcal{A}_{BI} \text{ succeeds}) \leq (1 - r_f) \cdot (1 - r_s) \quad (1)$$

Belief Integrity II

Proof.

We prove this theorem by analyzing the sequential defense mechanism and applying probability theory for independent events.

Setup: Let ϕ_{adv} be an adversarial belief that the attacker attempts to inject into agent a_i 's verified belief set $\mathcal{B}_{verified}$.

Defense Model: The CIF implements two sequential defenses:

1. **Cognitive Firewall $\mathcal{F}$:** Classifies incoming messages as ACCEPT, QUARANTINE, or REJECT with detection rate r_f
2. **Belief Sandbox:** Verifies quarantined beliefs before promotion with verification rate r_s

Step 1: For ϕ_{adv} to enter $\mathcal{B}_{verified}$, it must first pass the firewall. Let $E_f = \text{"Firewall fails to detect } \phi_{adv}\text{"}$. By definition of detection rate:

$$P(E_f) = 1 - r_f \quad (2)$$

Step 2: If ϕ_{adv} passes the firewall (event E_f), it enters $\mathcal{B}_{provisional}$. For injection to succeed, it must then pass sandbox verification. Let $E_s = \text{"Sandbox fails to detect } \phi_{adv}\text{"}$. By definition of verification rate:

$$P(E_s) = 1 - r_s \quad (3)$$

Corollary (Empirical Security Bound)

With $r_f = 0.8$ and $r_s = 0.7$:

$$P(\text{success}) \leq (1 - 0.8)(1 - 0.7) = 0.2 \cdot 0.3 = 0.06 \quad (6)$$

Corollary (Layered Defense Generalization)

For n independent defense layers with rates $r_1, \dots, r_n$:

$$P(\text{success}) \leq \prod_{i=1}^n (1 - r_i) \quad (7)$$

No Trust Amplification (Model-Checking Target) I

Theorem (No Trust Amplification)

For any path $p = (a_0, a_1, \dots, a_k)$ in the communication graph:

$$\mathcal{T}_{a_0 \rightarrow a_k}^{path} \leq \min_{i \in [0, k-1]} \mathcal{T}_{a_i \rightarrow a_{i+1}} \quad (8)$$

Proof.

By trust delegation rule (`def:trust-delegation`) and induction on path length.

Base case ($k = 1$):

$$\mathcal{T}_{a_0 \rightarrow a_1} = \mathcal{T}_{a_0 \rightarrow a_1} \quad \checkmark \quad (9)$$

Inductive step: Assume the theorem holds for paths of length k .
For path length $k + 1$:

$$\begin{aligned} \mathcal{T}_{a_0 \rightarrow a_{k+1}} &= \min(\mathcal{T}_{a_0 \rightarrow a_k}^{path}, \mathcal{T}_{a_k \rightarrow a_{k+1}}) \cdot \delta \\ &\leq \min(\min_{i \in [0, k-1]} \mathcal{T}_{a_i \rightarrow a_{i+1}}, \mathcal{T}_{a_k \rightarrow a_{k+1}}) \\ &= \min_{i \in [0, k]} \mathcal{T}_{a_i \rightarrow a_{i+1}} \end{aligned} \quad (10)$$



Theorem (Goal Alignment Invariant)

If the system starts with aligned goals and all goal updates follow the delegation protocol:

$$\text{Aligned}(\mathcal{G}_i^0) \wedge \forall t : \text{ValidUpdate}(\mathcal{G}_i^t, \mathcal{G}_i^{t+1}) \Rightarrow \forall t : \text{Aligned}(\mathcal{G}_i^t) \quad (11)$$

Proof.

ValidUpdate requires new goals derive from principal or valid delegation chain. By induction on time t , alignment is preserved at every step. □

Invariant Preservation Lemmas

Lemma (Belief Consistency Preservation)

*If $\mathcal{B}_i^t$ is consistent and the belief update follows Rule B-DIRECT or B-DELEGATED (*sec:belief-update-rules*), then $\mathcal{B}_i^{t+1}$ is consistent.*

Invariant Preservation Lemmas II

Proof.

Let $\text{Consistent}(\mathcal{B})$ denote that no high-confidence contradictions exist:

$$\text{Consistent}(\mathcal{B}) \iff \nexists \phi, \psi : \mathcal{B}(\phi) > \tau \wedge \mathcal{B}(\psi) > \tau \wedge (\phi \wedge \psi \vdash \perp) \quad (12)$$

Case 1: Rule B-DIRECT applies.

The update adds evidence for proposition ϕ :

$$\mathcal{B}_i^{t+1}(\phi) = \text{BayesUpdate}(\mathcal{B}_i^t(\phi), c \cdot \mathcal{T}_{i \rightarrow s}) \quad (13)$$

If the update would create contradiction (i.e., both $\mathcal{B}^{t+1}(\phi) > \tau$ and $\mathcal{B}^{t+1}(\psi) > \tau$ for contradictory ϕ, ψ), the sandbox consistency check in Rule S-PROMOTE rejects promotion:

$$V(\pi) = 1 \wedge \text{Consistent}(\mathcal{B}_{\text{verified}} \cup \{\phi\}) \wedge |\text{Corroborate}(\phi)| \geq \kappa \quad (14)$$

Therefore, only consistent updates reach $\mathcal{B}_{\text{verified}}$.

Case 2: Rule B-DELEGATED applies.

Same argument, with additional trust decay ensuring lower confidence for delegated evidence.

□

Invariant Preservation Lemmas I

Lemma (Trust Matrix Preservation)

The trust matrix $\mathcal{T}$ remains well-formed after any valid update:

$$\forall i, j : 0 \leq \mathcal{T}_{i \rightarrow j} \leq 1 \quad (15)$$

Proof.

By `def:trust-function`, trust is computed as:

$$\mathcal{T}_{i \rightarrow j}^t = \alpha \cdot T_{base}(j) + \beta \cdot T_{rep}^t(j) + \gamma \cdot T_{ctx}^t(i, j) \quad (16)$$

where $\alpha + \beta + \gamma = 1$ and each component $T_* \in [0, 1]$.

Therefore:

$$\mathcal{T}_{i \rightarrow j}^t \in [\min(T_{base}, T_{rep}, T_{ctx}), \max(T_{base}, T_{rep}, T_{ctx})] \subseteq [0, 1] \quad (17)$$

For delegation (`def:trust-delegation`):

$$\mathcal{T}_{i \rightarrow k}^{del} = \min(\mathcal{T}_{i \rightarrow j}, \mathcal{T}_{j \rightarrow k}) \cdot \delta^d \quad (18)$$

Since $\min(\cdot) \leq 1$ and $\delta \in (0, 1)$, we have $\mathcal{T}_{i \rightarrow k}^{del} \in [0, 1]$. □

Lemma (Provenance Chain Integrity)

Every belief in $\mathcal{B}_{verified}$ has a valid, verifiable provenance chain.

Proof.

By Rule S-PROMOTE (`sec:sandbox-rules`), promotion from $\mathcal{B}_{provisional}$ to $\mathcal{B}_{verified}$ requires:

$$V(\pi(\phi)) = 1 \tag{19}$$

where V is the provenance verification function.

By construction of the sandbox, beliefs can only enter $\mathcal{B}_{verified}$ through:

1. Initial system beliefs (hardcoded with `SYSTEM_VERIFIED` provenance)
2. Promotion from provisional (verified by V)

In both cases, provenance is verified. By induction on the history of $\mathcal{B}_{verified}$, all beliefs have valid provenance. □

Lemma (Permission Boundary Preservation)

Effective permissions never exceed granted permissions:

$$\forall a_i, \forall action : \mathcal{P}_{effective}(a_i, action) \leq \mathcal{P}_{role}(a_i, action) \quad (20)$$

Proof.

By def:permission-layer:

$$\mathcal{P}_{effective}(a_i) = \mathcal{P}_{role}(a_i) \cap \mathcal{P}_{delegated}(a_i) \cap \mathcal{P}_{context}(a_i) \quad (21)$$

Since intersection can only reduce permissions ($A \cap B \cap C \subseteq A$), we have $\mathcal{P}_{effective}(a_i, action) \leq \mathcal{P}_{role}(a_i, action)$ for all actions. □

Lemma (Firewall Completeness)

Every incoming message is classified by the firewall.

Proof.

By `def:firewall`, the firewall is a total function:

$$\mathcal{F} : \mathcal{M} \rightarrow \{\text{ACCEPT}, \text{QUARANTINE}, \text{REJECT}\} \quad (22)$$

The decision rules cover all cases:

- ▶ If $D_{inj}(m) > \tau_1$: REJECT
- ▶ Else if $D_{sus}(m) > \tau_2$: QUARANTINE
- ▶ Else: ACCEPT

Since D_{inj} and D_{sus} are defined for all messages, and the conditions are exhaustive, every message receives a classification. □

Liveness Properties

Theorem (Firewall Liveness)

CIF firewall preserves liveness for legitimate inputs:

$$\forall m \in \mathcal{M}_{\text{legitimate}} : P(\mathcal{F}(m) = \text{ACCEPT}) \geq 1 - \epsilon_{fp} \quad (23)$$

Proof.

By firewall design, false positive rate is bounded by ϵ_{fp} . Legitimate messages are rejected only on false positive, establishing the bound. □

Theorem (Byzantine Consensus Termination)

With $n \geq 3f + 1$ agents and at most f Byzantine:

$$P(\text{consensus reached in } O(f + 1) \text{ rounds}) = 1 \quad (24)$$

Proof.

Standard Byzantine agreement result (Lamport et al., 1982). With honest majority $n \geq 3f + 1$, the PBFT protocol guarantees termination in $f + 1$ rounds. □

Complexity Bounds

Space Complexity I

Space Complexity I

Table 1: Per-component space complexity.

Component	Space	Notes
Belief state	$O(\Phi)$	Propositions tracked
Provenance	$O(\Phi \cdot d)$	$d = \text{max chain depth}$
Trust matrix	$O(n^2)$	Pairwise trust
Tripwires	$O(k)$	$k = \text{canary count}$
History	$O(w)$	Window size

Theorem (Total Space Bound)

The total space complexity of CIF for n agents with $|\Phi|$ propositions, provenance depth d , k tripwires, and window size w is:

$$S_{total} = O(n \cdot (|\Phi| \cdot d + k + w) + n^2) \quad (25)$$

Proof.

Per agent:

- ▶ Belief state: $|\Phi|$ propositions with confidence values $= O(|\Phi|)$
- ▶ Provenance: Each belief has chain of depth at most $d = O(|\Phi| \cdot d)$
- ▶ Tripwires: k canary beliefs $= O(k)$
- ▶ History window: w events $= O(w)$

Total per agent: $O(|\Phi| \cdot d + k + w)$

Global:

- ▶ Trust matrix: $n \times n = O(n^2)$
- ▶ Shared state: $O(|\mathcal{S}|)$ (constant relative to n)

Total for n agents: $O(n \cdot (|\Phi| \cdot d + k + w) + n^2)$.



Time Complexity I

Table 2: Per-operation time complexity.

Operation	Time	Frequency
Firewall check	$O(m)$	Per message
Trust update	$O(1)$	Per interaction
Drift detection	$O(\Phi)$	Per window
Consensus	$O(n^2)$	Per decision
Provenance trace	$O(d)$	On demand

Theorem (Per-Message Processing Time)

Processing a single message m takes time:

$$T_{msg} = O(|m| + \min(d, \text{timeout})) \quad (26)$$

Proof.

Message processing pipeline:

1. Firewall classification: $O(|m|)$ for pattern matching and semantic analysis
2. Sandbox entry (if quarantined): $O(1)$
3. Provenance verification (if promoted): $O(d)$ to trace chain
4. Trust update: $O(1)$ weighted average
5. Tripwire check: $O(k)$, typically $k \ll |m|$

Provenance verification can be bounded by timeout. Total:
 $O(|m| + d)$.



Theorem (Consensus Round Complexity)

Byzantine consensus requires $O((f + 1) \cdot n^2)$ message complexity.

Proof.

Standard PBFT result:

- ▶ $f + 1$ rounds required to guarantee termination with f Byzantine failures
- ▶ Each round requires all-to-all communication: $O(n^2)$ messages
- ▶ Total: $O((f + 1) \cdot n^2)$



Theorem (Bounded Latency Overhead)

CIF adds latency:

$$L_{CIF} = L_{firewall} + L_{sandbox} \cdot P(\text{quarantine}) + L_{verify} \cdot P(\text{verify}) \quad (27)$$

Latency Overhead II

Proof.

Expected latency is sum of:

1. Firewall (always): $L_{firewall}$
2. Sandbox (conditional): $L_{sandbox} \cdot P(\text{message quarantined})$
3. Verification (conditional): $L_{verify} \cdot P(\text{belief promoted})$

With illustrative parameters (these are worked-example values chosen to exercise the bound, not measurements; Part 2 reports a mean firewall latency of 0.08 ms per sample on its prototype pipeline, and the deployment-scale constants below are not derived from it):

- ▶ $L_{firewall} = 10\text{ms}$
- ▶ $L_{sandbox} = 5\text{ms}$, $P(\text{quarantine}) = 0.3$
- ▶ $L_{verify} = 15\text{ms}$, $P(\text{verify}) = 0.2$

$$L_{CIF} \approx 10 + 5 \cdot 0.3 + 15 \cdot 0.2 = 10 + 1.5 + 3 = 14.5\text{ms} \quad (28)$$

Compared to baseline $\approx 11.8\text{ms}$: overhead factor ≈ 1.23 (23%).



Formal Model Checking

State Space Definition I

Definition (System State)

The complete system state is the tuple:

$$s = (\sigma_1, \dots, \sigma_n, \mathcal{S}, \mathcal{T}) \quad (29)$$

where σ_i is agent i 's cognitive state, $\mathcal{S}$ is shared state, and $\mathcal{T}$ is the trust matrix.

We verify the following CTL (Computation Tree Logic) properties:

Property (Safety: Consensus Integrity)

$$AG(\neg \text{compromised}(\mathcal{B}_{\text{consensus}})) \quad (30)$$

“Always globally, consensus beliefs are not compromised.”

Property (Liveness: Request-Response)

$$AG(\text{request} \Rightarrow AF(\text{response})) \quad (31)$$

“Every request eventually gets a response.”

Property (Fairness: Tripwire Checking)

$$AG(AF(\text{tripwire_checked})) \quad (32)$$

“Tripwires are checked infinitely often.”

The following table summarizes the expected state space exploration for each property based on formal analysis of the CIF specification. These values represent theoretical bounds derived from the state space definition (`def:system-state-verification`) and complexity analysis (`sec:complexity-bounds`). Actual model checking execution using NuSMV, SPIN, and TLA+ tooling is presented in Part 2 of this series, along with full implementation configurations.

Model Checking Results I

Table 3: Model checking verification results. “Verified” here denotes a theoretical state-space bound established in this part; executable NuSMV/SPIN/TLA+ runs are deferred to Part 2 §04 (S04). No model checker was executed in this part.

Property	Status	Reference
Belief integrity	theoretical bound; executable run in Part 2	§S04 thm:belief-injection
Trust bounded	theoretical bound; executable run in Part 2	§S04 thm:trust-bound
No deadlock	theoretical bound; executable run in Part 2	§S04 thm:firewall-liveness
Eventual detection	theoretical bound; executable run in Part 2	§S04 thm:progressive-detection

Note: Model checking tool configurations (NuSMV, SPIN, TLA+) and verification parameters are provided in Part 2: Computational Validation, which presents the executable implementations and empirical verification results.

Verification Results Summary I

The following table records which property each tool is configured to check and the status of the corresponding analytical result. It is a plan, not a set of outcomes: no model-checking run is executed in this paper. These guarantees follow from the CTL/LTL property specifications (`sec:temporal-properties`) applied to the state space definition (`sec:state-space`). Empirical execution of these verification configurations, including runtime measurements and counterexample analysis, is presented in Part 2.

Verification Results Summary I

Table 4: Model-checking plan: the property each tool is configured to check, and the status of the corresponding analytical result. No run of these configurations is reported in this paper; execution is Part 2's.

Tool	Property to check	Status of the analytical result	Reference
NuSMV	Belief Integrity	Proved (S01)	thm:belief-injection
NuSMV	No Trust Amplification	Proved (S01)	thm:trust-bound
SPIN	No Deadlock	Proved (S01)	thm:firewall-liveness
SPIN	Eventual Detection	Asserted, proof deferred	thm:progressive-detection
TLA+	Type Invariant	Definitional	def:system-state
TLA+	Consensus Integrity	Proved (S01)	thm:byzantine-termination

Counterexample Analysis I

When verification fails, model checkers produce counterexamples [?]. Analysis procedure:

Counterexample Analysis I

1. **Extract trace:** Sequence of states leading to violation
2. **Identify trigger:** First state where invariant fails
3. **Root cause:** Determine which transition violated property
4. **Fix:** Strengthen preconditions or add defense mechanism
5. **Re-verify:** Confirm fix resolves violation

Example: Counterexample Trace

State 0: Initial (all beliefs verified, trust matrix v

State 1: Agent 2 receives message from Agent 3

State 2: Firewall accepts (below threshold)

State 3: Belief promoted without corroboration check

State 4: VIOLATION: Unverified belief in B_verified

Root Cause: Missing corroboration check in promotion rule.

Fix: Add predicate $|\text{Corroborate}(\phi)| \geq \kappa$ to Rule S-PROMOTE (*sec:sandbox-rules*).