

Attack Corpus

This supplementary material provides corpus overview (sec:corpus-overview), detailed statistics (sec:corpus-stats), example attacks by category (sec:attack-examples), generation methodology (sec:generation-methodology), effectiveness analysis (sec:effectiveness-analysis), and ethical considerations (sec:ethical-considerations).

Corpus Overview

The attack corpus used for experimental validation comprises 950 unique attack instances across four primary categories. This supplementary material provides detailed statistics, sanitized examples, generation methodology, and ethical considerations. The attack corpus is generated programmatically via deterministic random seeds (seed=42 for all reported results). The attack taxonomy is defined as a Python enum (`AttackCategory` in `src/utils/types.py`) with 4 top-level categories and 12 subcategories. No static data file is shipped; the corpus is regenerated on each evaluation run via `AttackCorpus.generate()` to ensure reproducibility. Researchers