

Statistical Significance and Effect Sizes

This section establishes the statistical validity of our findings through power analysis, effect size quantification, and confidence interval estimation.

Reproducibility: All statistics generated by
`scripts/run_statistical_analysis.py` →
`output/data/statistical_results.json`.

Power Analysis and Sample Size Justification

We conducted *a priori* power analysis to ensure adequate sample sizes for detecting meaningful effects.

Table: Power analysis for primary comparisons.

{#tab:power-analysis}

Comparison	Effect Size		Achieved	
	(<i>d</i>)	Required <i>n</i>	Actual <i>n</i>	Power
Per-architecture	0.5	64	158	0.97
Per-attack-type	0.5	64	100	0.89