

The Frontier: Where We Need You

We have built the foundation and validated the core mechanics. However, the field is in its infancy. We have established valid trust metrics (Trust Calculus) and filtering mechanisms (Firewalls), but we lack standardized protocols for agent state persistence and identity federation—the agentic equivalents of cookies or OAuth.

This is where you come in.

The Cognitive Integrity Framework is open source, and the problems below are not just “future research”—they are immediate engineering blockers that need to be solved.

1. The UX of Trust (Help Wanted)

The Problem: Currently, when an agent flag is raised (“Identity Violation”), it looks like a JSON error log. **The Need:** We need a “Cognitive Dashboard” for human operators. What does it look like to visualize the trust graph of 50 agents in real-time? How do we show “Drift” intuitively? **The Goal:** A React/Next.js dashboard that connects to the CIF Python SDK.

Contributing

This is not a closed academic project. It is a living defense framework for the agentic future.

- ▶ **Code:** github.com/docxology/cognitive_integrity
- ▶ **Discussion:** Join the `discussions` tab on GitHub.
- ▶ **Contribute:** We prioritize PRs that add **Adapters** for new agent frameworks (e.g., Semantic Kernel, Microsoft AutoGen).

Let's build the immune system for the agentic web, together.