

Where We Stand: A Call to Build

This series began with a theory (Part 1) and moved to an experiment (Part 2). It ends here, with a call to engineering.

The Theory Holds

We proved that trust can be bounded. We proved that defenses can be composed algebraically. We proved that stealth and impact are inversely related. These are not just academic curiosities; they are foundational constraints for secure cognitive systems.

The Code Works

We validated these laws in code. 1,594 tests. 950 attacks. 97% structural constraint satisfaction. The system works. It is not hypothetical.

The Next Step is Yours

As we move beyond simple prompt engineering, the era of **Cognitive Security Engineering** has begun.

We are no longer just asking LLMs to write poems; we are asking them to run the world. If we want them to do that safely, we cannot rely on luck or better fine-tuning. We need structure. We