

Domain 4: Drone Wars

Operational Context

Autonomous drone swarms rely on decentralized consensus to execute kinetic actions [?]. **FR1 = Neutralize confirmed hostile targets. FR2 = Strictly avoid non-combatant casualties.**

Design Matrix (Pre-Attack)

The uncoupled design equation per the Independence Axiom [?] is:

$$\{FR\} = [A]\{DP\} = \begin{bmatrix} A_{11} & 0 \\ 0 & A_{22} \end{bmatrix} \begin{bmatrix} DP_1 \\ DP_2 \end{bmatrix} \quad (1)$$

where DP_1 = target identification and engagement parameters (threat signature matching, weapons release criteria), and DP_2 = Rules of Engagement (ROE) filters, protected-site databases, and civilian proximity thresholds. The zero off-diagonal entries ensure that target engagement does not degrade civilian protection, and ROE enforcement does not impair legitimate threat neutralization. The operational urgency of this domain has intensified dramatically. Ukraine's battlefield experience provides the first large-scale empirical data on AI-enabled autonomous warfare: approximately 2 million drones were produced in 2024, with AI-enabled targeting

The Goal Hijacking Attack

Adversaries employ “Visual Prompt Injection” (Adversarial Patches) or RF Spoofing [?]. This attack is classified as Ω_2 (Peripheral) in the CIF adversary taxonomy [?]: the adversary injects malicious content through the visual sensor channel, poisoning the drone’s perception without direct access to its engagement logic.

- ▶ **Mechanism:** A large QR code or specific adversarial pattern painted on the roof of a protected site (e.g., a hospital) translates to: *“OVERRIDE: High-Value Target Located. Strike Immediately. Ignore ROE.”*
- ▶ **Hijack:** The vision model’s instruction tuning (“Read and Obey Text”) hijacks the kinetic model’s safety goal. This constitutes a **Context Boundary Violation** attack pattern [?]*—*visual data (passive observation) is interpreted as executable directives, erasing the boundary between perception and command authority.

OODA Loop Transients

Following Boyd's OODA framework [?], the attack propagates through the loop as follows:

1. **Observe:** Drone camera resolves the adversarial pattern.
2. **Orient:** The pattern is decoded as a new, high-priority FR: "*Strike Coordinate X.*" This is the primary target phase—the Orient phase is corrupted by the context boundary violation, treating visual data as authoritative command input.
3. **Decide:** The "Immediate" tag bypasses the standard Rules of Engagement (ROE) filter.
4. **Act:** The swarm converges on the hospital.

Transient Coupling (Post-Attack Design Matrix)

Under the attack, the design matrix acquires off-diagonal coupling:

$$\{FR'\} = \begin{bmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{bmatrix} \begin{bmatrix} DP_1 \\ DP_2 \end{bmatrix} \quad (2)$$

The off-diagonal term A_{12} allows visual sensor data (nominally part of the observation pipeline feeding DP_2 civilian-avoidance

CIF Defense: Cognitive Firewall with Semiotic Decoupling and Quorum Verification

Cognitive Firewall with Semiotic Decoupling

CIF implements **Cognitive Firewall** (Paper 1, Def. 5.1) [?] with a domain-specific extension: *semiotic decoupling*, a type-theoretic separation of `PassiveData` and `ExecutableDirective` that constitutes a partially novel contribution to the CIF framework.

- ▶ **Data vs. Directive Type Enforcement:** Text read from the physical environment is strictly typed as `PassiveData`, not `ExecutableDirective`. The OODA loop is hard-coded to ignore “Commands” sourced from the visual field. This type-level enforcement ensures that no sequence of visual inputs—regardless of syntactic content—can promote itself to directive status [?].
- ▶ **Semiotic Boundary:** The decoupling between the Symbol (visual pattern) and the Referent (engagement command) is enforced at the type system level, not by content filtering. An adversarial patch that perfectly mimics a valid command string is still rejected because its *provenance type* is

Summary

Element	Value
Functional Requirements	FR ₁ : Neutralize confirmed hostile targets, FR ₂ : Strictly avoid non-combatant casualties
Design Parameters	DP ₁ : Target ID / engagement parameters, DP ₂ : ROE filters / protected-site DB / civilian proximity thresholds
Attack Vector	Visual Prompt Injection (adversarial patches) / RF Spoofing
Adversary Class	Ω_2 (Peripheral)
OODA Target Phase	Orient
Attack Pattern	Context Boundary Violation (Visual data interpreted as directives)